

GPU Computing: A Quick Start

Orest Shardt

Department of Chemical and Materials Engineering
University of Alberta

August 25, 2011

Session Goals

- Get you started with highly parallel LBM
- Take a practical approach
 - Less about hardware/architecture details
 - More about how to use it
 - Show performance impact of programming decisions
- Focus on what is possible with GPUs
- Detailed reference information is available on manufacturers' websites

CPUs

- Sequential operation (on each core)
 - Read memory
 - Compute
 - Store to memory
- Sometimes, need to perform exactly the **same operation** but on **different data**
- Order of execution does not matter

Pseudo-code for applying collision operator

```
for(int i = 0; i < NX; ++i)
  for(int j = 0; j < NY; ++j)
    collide(distribution[i][j]);
```



```
collide(distribution[0][0]);
collide(distribution[0][1]);
```

...

```
collide(distribution[NX-1][NY-2]);
collide(distribution[NX-1][NY-1]);
```

GPUs for Graphics



- For 3D graphics, need to process large meshes very quickly
- You perform the **same operation** (determine visibility, compute normal vector, shade) on a lot of **different data** (triangles)
- Led to development of highly specialized “stream” processors
- Architecture was adapted for general purpose use: GPGPU

GPUs for Computing

- Parallel operation
 - Read data set from memory
 - Perform the **same computation** on many data sets **simultaneously**
 - Store results to memory

Pseudo-code for collision operator

```
for(int i = 0; i < NX; ++i)
  for(int j = 0; j < NY; ++j)
    collide(distribution[i][j]);
```





N locations processed simultaneously

collide(distribution[0][0]);	—	collide(distribution[0][1]);	—	collide(distribution[0][N-1]);
collide(distribution[0][N]);	—	collide(distribution[0][N+1]);	—	collide(distribution[0][2N-1]);
...	
collide(distribution[1][0]);	—	collide(distribution[1][1]);	—	collide(distribution[1][N-1]);
collide(distribution[1][N]);	—	collide(distribution[1][N+1]);	—	collide(distribution[1][2N-1]);
...	
collide(distribution[NX-1][0]);	—	collide(distribution[NX-1][1]);	—	collide(distribution[NX-1][N-1]);
...	
collide(distribution[NX-1][NY-N]);	—	collide(distribution[NX-1][NY-N+1]);	—	collide(distribution[NX-1][NY-1]);

GPU Performance


- When a problem can be parallelized, the execution speed can increase dramatically
- High speed at low cost
 - GTX 580 ~\$500, Intel i7 Quad Core ~\$300



From NVIDIA CUDA C Programming Guide


Available Architectures

- AMD/ATI
 - Radeon
- NVIDIA
 - GeForce, Tesla, Quadro



Development Environments

- Architecture Specific
 - Make use of specific hardware capabilities
 - NVIDIA: Compute Unified Device Architecture (CUDA)
 - AMD/ATI: AMD Accelerated Parallel Processing (APP) SDK (formerly ATI Stream)
- General Purpose
 - Work with NVIDIA and AMD GPUs
 - OpenCL
 - standard for heterogeneous computing: CPU, FPU, GPU, embedded, etc
 - DirectCompute
 - part of Direct X for Windows




Languages

- C/C++ is best supported
- CUDA Fortran is available (Portland Group)
- OpenCL bindings are available for many languages
 - Java, Javascript, Python, C#
- GPGPU capabilities are available in other computing packages, e.g. Matlab, Mathematica

We will work with

- NVIDIA CUDA Toolkit 4.0 (May 2011)
- C/C++





EXAMPLE

Your first kernel

- CUDA C/C++
 - All features of C/C++ are allowed in code that executes on the “host” cpu
 - Some restrictions on “device” (gpu) code
 - *.cu* source files combine host and device code

vectorAdd.cu

Regular includes
for standard
libraries

```
#include <stdlib.h>
#include <stdio.h>
```

Identifies this
function as device
code that is called
by host

```
// the kernel
__global__ void vecAddKernel(float* a, float* b, float* c)
{
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x;
    c[i] = a[i]+b[i];
}
```

Special variables
available in
device code

vectorAdd.cu (continued)

no `__global__`
means this is
host code


```
int main(int argc, char** argv)
{
    // use command-line specified CUDA device, otherwise use
    // device with highest Gflops/s
    if(cutCheckCmdLineFlag(argc, (const char**)argv, "device"))
        cutilDeviceInit(argc, argv);
    else
        cudaSetDevice(cutGetMaxGflopsDeviceId());

    int deviceId;
    // get information about the selected device
    cudaGetDevice(&deviceId);
    cudaDeviceProp deviceProp;
    cudaGetDeviceProperties(&deviceProp, deviceId);

    // constants for vector size and memory size
    const unsigned int vector_size = 12*1024*1024;
    const unsigned int mem_size = sizeof(float)*vector_size;

    printf("Adding %i MB vectors\n", vector_size/(1024*1024));
    printf("Using device %i: %s\n", deviceId, deviceProp.name);
```

CUDA API
functions



vectorAdd.cu (continued)

```
// allocate host memory
float* a_host = (float*) malloc(mem_size);
float* b_host = (float*) malloc(mem_size);
float* c_host = (float*) malloc(mem_size);

// initialize the memory
for(unsigned int i = 0; i < vector_size; ++i)
{
    a_host[i] = 1.0f*i;
    b_host[i] = 2.0f*i;
    c_host[i] = 0.0f;
}

// allocate device memory
float* a_dev;
float* b_dev;
float* c_dev;
cudaMalloc((void**) &a_dev, mem_size);
cudaMalloc((void**) &b_dev, mem_size);
cudaMalloc((void**) &c_dev, mem_size);
```

vectorAdd.cu (continued)

```
// create events for timing
cudaEvent_t start, stop;
float time,gbps;
cudaEventCreate(&start);
cudaEventCreate(&stop);

// copy from host to device memory
cudaMemcpy(a_dev, a_host, mem_size, cudaMemcpyHostToDevice);
cudaMemcpy(b_dev, b_host, mem_size, cudaMemcpyHostToDevice);
cudaMemcpy(c_dev, c_host, mem_size, cudaMemcpyHostToDevice);

// get ready to execute kernel

// kernel execution parameters
unsigned int nThreads = 512;
unsigned int nBlocks = vector_size/nThreads;

// blocks in grid
dim3 grid(nBlocks, 1, 1);
// threads in block
dim3 threads(nThreads, 1, 1);
```

vectorAdd.cu (continued)

```
cudaEventRecord(start,0);           // get time before kernel starts

// execute the kernel
vecAddKernel<<< grid, threads >>>>(a_dev, b_dev, c_dev);

cudaEventRecord(stop,0);           // get time after kernel finishes
cudaEventSynchronize(stop);        // wait for GPU to finish
cudaEventElapsedTime(&time,start,stop); // compute time difference

printf("runtime: %.4f (ms)\n", time);

// copy result from device to host
cudaMemcpy(c_host, c_dev, mem_size, cudaMemcpyDeviceToHost);
```

triple angle bracket syntax for launching a kernel

vectorAdd.cu (continued)

```
// release event resources
cudaEventDestroy(start);
cudaEventDestroy(stop);

// free host memory
free(a_host);
free(b_host);
free(c_host);

// free device memory
cudaFree(a_dev);
cudaFree(b_dev);
cudaFree(c_dev);

cudaDeviceReset();

} // end of main()
```

Sample Results for Full Program

Compiling and running:

```
nvcc --ptxas-options=-v -arch sm_13 -O3 -c -I $(CUDADIR)/C/common/inc -o vecAdd.cu.o vecAdd.cu
g++ -c -O3 vecAdd.cpp -o vecAdd.cpp.o
g++ vecAdd.cu.o vecAdd.cpp.o -o vecAdd -lcudart -L$(CUDALIBDIR) $(CUDADIR)/C/lib/libcutil_x86_64.a
./vecAdd
```

Adding 12 MB vectors


Using device 0: GeForce 310M

Host to device memory copy: 48.0MB in 97.6 (ms)
Bandwidth: 1.44 (GB/s)

vecAddKernel

runtime: 24.6969 (ms)
bandwidth: 5.69 (GB/s)
Correct result: YES

Device to host memory copy: 48.0MB in 35.1 (ms)
Bandwidth: 1.34 (GB/s)




ARCHITECTURE


NVIDIA Architecture

- Stream processors
 - SIMT: Single Instruction, Multiple Thread
 - Similar to SIMD: Single Instruction, Multiple Data

Schematics of CPU and GPU organization



CPU




GPU

Threads, Blocks, and Grids

- **Kernels** are special functions that are executed simultaneously by many threads
- **Threads** are organized in a 3D arrangement called a **block**
- **Blocks** are organized in a 3D **grid**
- Blocks must be able to execute independently and in any order
- Blocks are split into **warps** for execution
- Special variables in device code
 - threadIdx, blockIdx, blockDim, gridDim

2D example of grid and block hierarchy




Block Execution

- Automatic parallelization
- Don't need to think about device details
- But, still need to know resource limits...


Serial Execution



2 Core GPU



4 Core GPU



Block and Grid Sizes

- Choosing block and grid sizes
 - Try to match problem
 - Use many threads to hide memory access delays
 - Device limits on block and grid dimensions
 - Check reference or run deviceQuery example
 - Memory limits
 - Registers: $threadsPerBlock * registersPerThread < maxRegistersPerBlock$
 - Shared and constant memory
 - compiler tells you: `nvcc --ptxas-options=-v`

Sample deviceQuery output

```
Total number of registers available per block: 16384
Maximum number of threads per block:          512
Maximum sizes of each dimension of a block:    512 x 512 x 64
Maximum sizes of each dimension of a grid:     65535 x 65535 x 1
```

Memory

- Local to a thread
 - registers (fastest)
 - managed by compiler
- Local to all threads in block
 - shared
 - used for cooperation between threads
 - managed by user
 - `__shared__`
- Visible to host and device
 - global (slowest)
 - managed by user



```
cudaError_t cudaMalloc(void **devPtr, size_t size)
cudaError_t cudaMemcpy(void *dst, const void *src, size_t count, enum cudaMemcpyKind kind)
cudaError_t cudaFree(void *devPtr)
```




PERFORMANCE

Timing

```
// create events for timing
cudaEvent_t start, stop;

cudaEventCreate(&start);
cudaEventCreate(&stop);


cudaEventRecord(start,0);
// do something interesting
cudaEventRecord(stop,0);
cudaEventSynchronize(stop);
float time;

// compute time difference
cudaEventElapsedTime(&time,start,stop);
// release event resources
cudaEventDestroy(start);
cudaEventDestroy(stop);
```



- Don't use CPU time
- If using wall time, `cudaDeviceSynchronize()`


Branching



- Lesson: structure code to avoid divergent execution paths in threads

Floating Point

- <1.3: only single precision (32 bit) floating point
- ≥ 1.3 : single and double precision (64 bit)



- IEEE 754 standard defines rules for floating point computation
- Due to rounding, $(A+B)+C \neq A+(B+C)$
- GPU provides a fused multiply add (FMA) instruction
 - round only once after $A*B+C$

Floating Point

- Compliance with IEEE 754 depends on device compute capability
- On early devices, some single precision operations do not have full single precision accuracy
- Now (capability >2.0), both single and double precision computations are IEEE compliant
- Lesson: Do not expect identical results on GPU and CPU
- Question: Which is better?
- For more information, see Whitehead and Fit-Florea, “Precision & Performance: Floating Point and IEEE 754 Compliance for NVIDIA GPUs”

Memory

- **Most important optimization** for memory intensive algorithm like LBM
- This is also important in CPUs
 - Lab: we will see that poor memory use causes $\sim 10\times$ slow down
- Don't compare poor CPU code with good GPU code or vice versa
- General guide:
 - **CPU: *consecutive*** memory accesses should be nearby
 - **GPU: *simultaneous*** memory accesses should be nearby
- Lesson: watch memory access patterns in both CPUs and GPUs

Bandwidth

$$\textit{Theoretical Bandwidth} = \textit{Memory Clock} \times \textit{Bus width (bits)} \times \frac{1 \text{ byte}}{8 \text{ bits}} \times 2$$


Divide by 10^9 or 1024^3 to get GB (use consistent definition)

Double Data Rate (DDR) memory:
2 transactions per clock cycle

Device 0: "GeForce GTX 460"

Memory Clock rate:

1800.00 Mhz

Memory Bus Width:

256-bit

$$\begin{aligned} \textit{Theoretical Bandwidth} &= 1.8 \times 10^9 \text{ s}^{-1} \times 256 \text{ bits} \times \frac{1 \text{ byte}}{8 \text{ bits}} \times \frac{1 \text{ GB}}{1024^3 \text{ bytes}} \times 2 \\ &= 107.3 \text{ GB/s} \end{aligned}$$


$$\textit{Effective Bandwidth} = \frac{\textit{bytes}_{read} + \textit{bytes}_{written}}{\textit{time interval}}$$

Coalesced Access


- 1.x: memory accesses in a half-warp (16 threads) are combined into one transaction if:
 - 1.0, 1.1: the i^{th} thread accesses the i^{th} word (32 bits) of an aligned 16 word (64 byte) segment
 - Otherwise, one transaction per thread
 - 1.2, 1.3: all threads access words (in any order) in an aligned 16 word segment
 - Threads may be combined to reduce number of transactions and issue smaller transactions
- 2.x: global memory access is cached
 - alignment and access patterns still matter, like on CPU

Memory Access Patterns

■ = one 32-bit word
■ = one thread




Result: one 64 byte transaction



Result


- 1.0, 1.1: 16 transactions
- >1.1: if aligned, one 128 byte transaction otherwise, one 64 byte and one 32 byte

Memory aligned to 64 bytes



Result

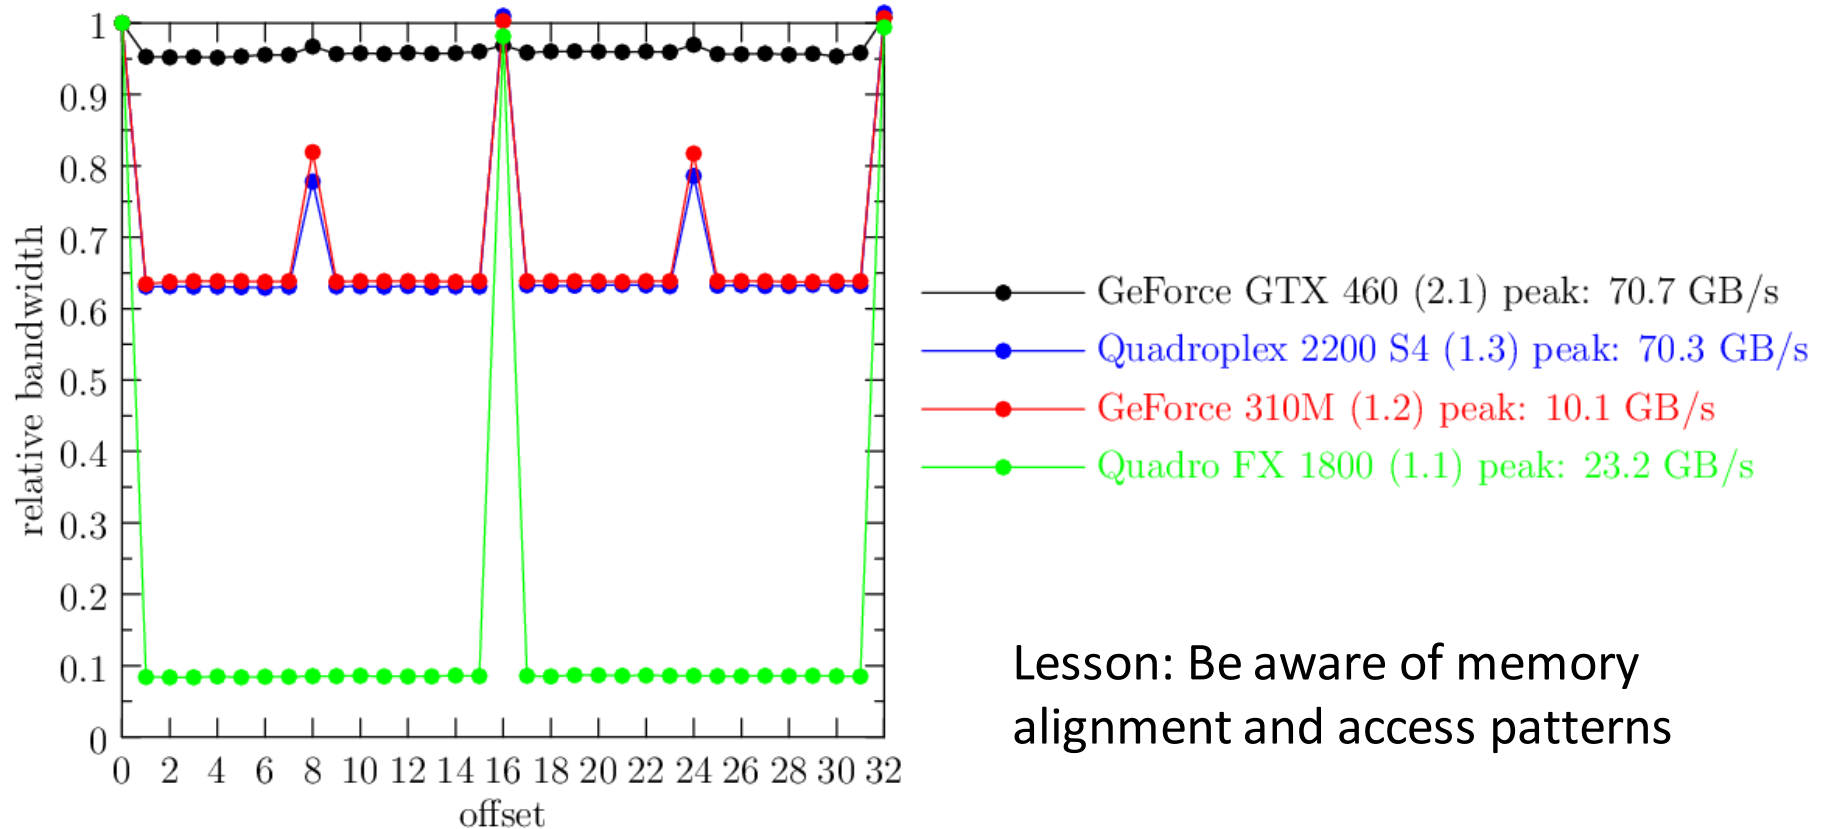
- 1.0, 1.1: 16 transactions
- >1.1: one 64 byte transaction



Result

- 1.0, 1.1: 16 transactions
- >1.1: if aligned, one 128 byte transaction otherwise, two 32 byte transactions


Non-Coalesced Access



Lesson: Be aware of memory alignment and access patterns

```
__global__ void readWriteKernel(float* a, int offset)
{
    unsigned int i = blockIdx.x*blockDim.x + threadIdx.x + offset;
    a[i] = 2.0f*a[i]+1.5f;
}
```

More Information





- Online

<http://developer.nvidia.com/nvidia-gpu-computing-documentation>

<http://developer.nvidia.com/cuda-education-training>

- Books and online tutorials are also available



LBM ON GPU

This afternoon


- Taylor-Green vortex decay using BGK
- Time dependent solution

$$u = u_{max} \sin(x) \cos(y) e^{-2\nu t}$$


$$v = -u_{max} \cos(x) \sin(y) e^{-2\nu t}$$

$$P = \frac{\rho}{4} (\cos(2x) + \cos(2y)) e^{-4\nu t}$$

- Domain: $0 \leq x, y \leq 2\pi$
- Implemented as a fully periodic domain
- 3 kernels
 - Initialization
 - Collision
 - Streaming



GPU Memory Layout

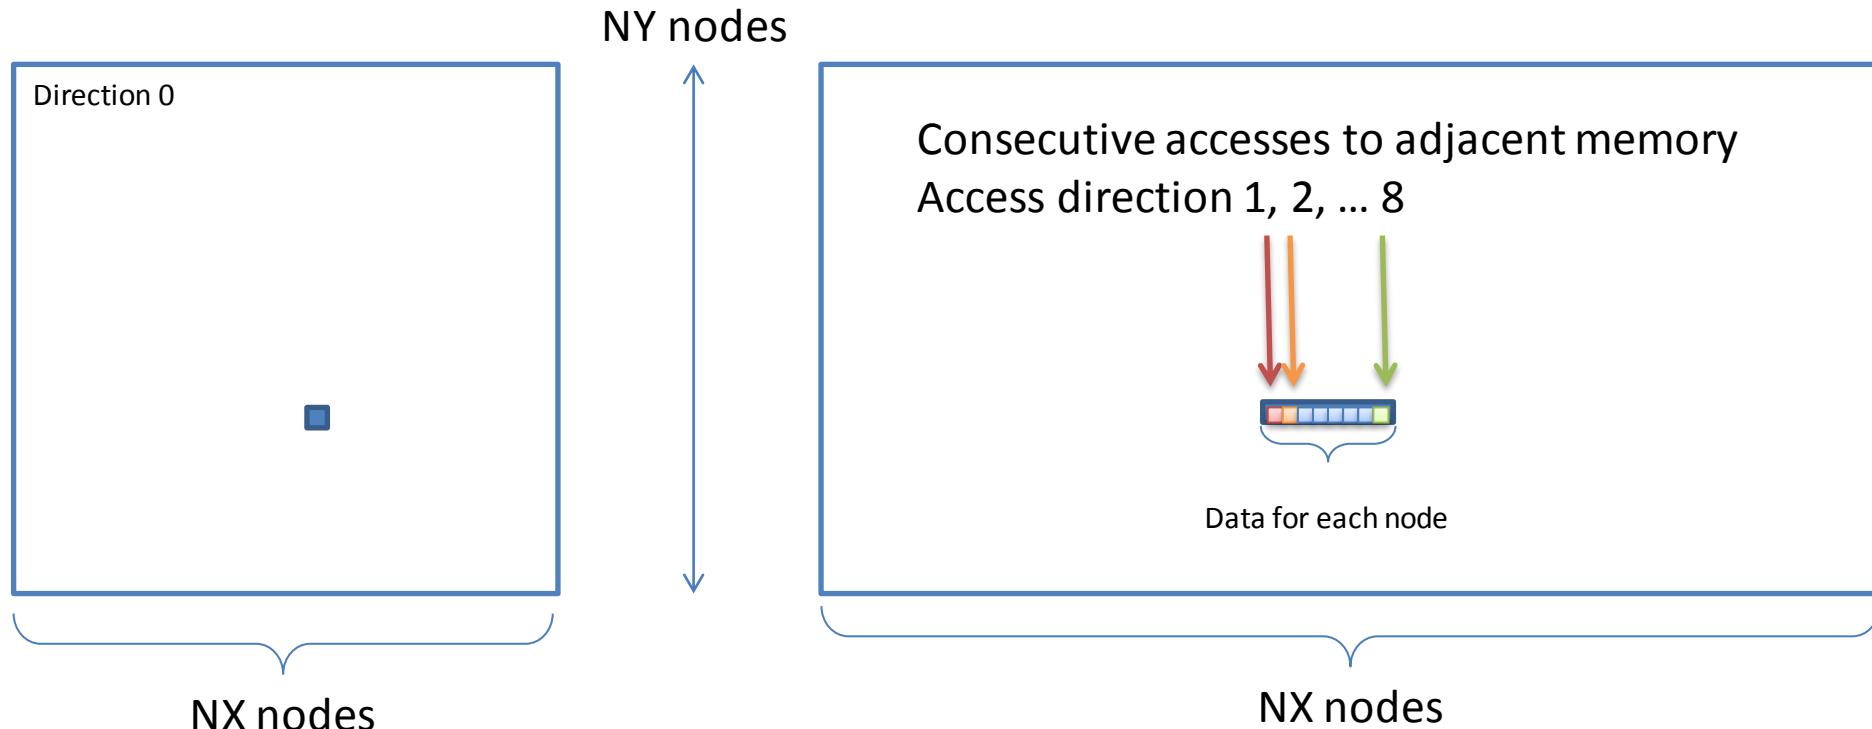


Thread i reads direction 0, then 1 ... 8
 Thread $i+1$ reads direction 0, then 1 ... 8
 Simultaneous accesses to adjacent memory

```
__global__ void gpu_collide(float* f0, float* f1)
{
    unsigned int y = blockIdx.y;
    unsigned int x = blockIdx.x*blockDim.x+threadIdx.x;
    ...
}
```

float* f0: separate array for direction 0
 avoids need for copying during streaming
float* f1: array for directions 1-8
float* f2: temporary array for streaming
 swapped with f1

CPU Memory Layout



Separate array for direction
0 to avoid need to copy
when streaming

```
extern "C" void cpu_collide(float* f0, float *f1)
{
    for(int y = 0; y < NY; ++y)
        for(int x = 0; x < NX; ++x)
        {
            ...
        }
}
```

LBM on GPU Performance

Simulating Taylor-Green Vortex Decay in a 1024x1024 2D domain for 8000 iterations with $u_{\max}=0.01$

Using device 0: **GeForce GTX 460**

Grid dimensions (32,1024,1)

Block dimensions (32,1,1)

Compute capability 2.1

GPU timing information

runtime: 20.6881 (s)

bandwidth: 51.36 (GB/s)

LBM speed: 405.48 (Mlups)

6.2 times faster than OpenMP parallelization for 6 core AMD Phenom II X6 1100T (3.3 GHz)

~37 times faster than a single core

Compare with 70.1 GB/s for memory access example

Device to host memory copy: 36.0MB in 32.8 (ms)

Bandwidth: 1.07 (GB/s)

CPU timing information

runtime: 128.8069 (s)

bandwidth: 8.25 (GB/s)

LBM speed: 65.13 (Mlups)

10.9 Mlups per core

Maximum absolute error between GPU and CPU solutions: $6.85453e-06$

Maximum relative error between GPU and CPU solutions: $3.87587e-05$

GPU L2 error: $2.4645e-07$

CPU L2 error: $1.04671e-07$

